

Local Search

Alice Gao

Lecture 6

Readings: R & N 4.1

Based on work by K. Leyton-Brown, K. Larson, and P. van Beek

Outline

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Hill climbing

- Hill climbing with random restarts

- Simulated Annealing

- Genetic Algorithms

Revisiting the Learning goals

Learning Goals

By the end of the lecture, you should be able to

- ▶ Describe the advantages of local search over other search algorithms.
- ▶ Formulate a real world problem as a local search problem.
- ▶ Given a local search problem, verify whether a state is a local/global optimum.
- ▶ Describe strategies for escaping local optima.
- ▶ Trace the execution of hill climbing, hill climbing with random restarts, and simulated annealing.
- ▶ Compare and contrast the properties of local search algorithms.

Why Use Local Search?

- ▶ Many search spaces are too big for systematic search.
- ▶ For CSPs, we only need to find a goal node. The path to a goal is irrelevant.
- ▶ Solution: local search

What is local search?

- ▶ Keep track of a single state, which is a complete assignment of values to variables.
- ▶ Move to a neighbour of the state based on how good the neighbour is.

When should we use local search?

- ▶ The state space is large or infinite.
- ▶ Memory is limited.
- ▶ To solve pure optimization problems.

Local Search

A local search problem consists of a:

- ▶ A **state** in the search space is a complete assignment to *all* of the variables.
- ▶ A **neighbour relation**: which states do I explore next?
- ▶ A **cost function**: how good is each state?

4-Queens Problem as a Local Search Problem

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Hill climbing

- Hill climbing with random restarts

- Simulated Annealing

- Genetic Algorithms

Revisiting the Learning goals

Questions to think about

The problem formulation:

- ▶ What is the neighbour relation?
- ▶ What is the cost function?

Executing the algorithm:

- ▶ Where do we start?
- ▶ Which neighbour do we move to?
- ▶ When do we stop?

Properties and performance of the algorithm:

- ▶ Given enough time, will the algorithm find the global optimum solution?
- ▶ How much memory does it require?
- ▶ How does the algorithm perform in practice?

Hill climbing

- ▶ Where do we start?
Start with a random solution.
- ▶ Which neighbour do we move to?
Move to a neighbour with the lowest cost. Break ties randomly.
- ▶ When do we stop?
Stop when no neighbour has a lower cost.
- ▶ How much memory does it require?
Only need to remember the current node.
No memory of where we've been.

Hill climbing in one sentence

Climbing Mount Everest in a thick fog with amnesia

CQ: Will hill climbing find the global optimum?

CQ: Given any problem and any initial state for the problem, will hill climbing always find the global optimum?


(A) Yes.

(B) No.

Trace the execution of hill climbing

See the **notes** on the course website.

The State Space Landscape



Where can hill climbing get stuck?

- ▶ Local optima: A state s^* is locally optimal if $c(s^*) \leq c(s)$ for every state s in the neighbourhood of s^* .
- ▶ Global optimum: A state s^* is globally optimal if $c(s^*) \leq c(s)$ for every state s .
- ▶ Plateau: a flat area in the state space
 - ▶ A shoulder: possible to improve when we escape.
 - ▶ A flat local optimum: cannot improve even if we escape.

CQ: Local and global optimum (1)

CQ: Consider the following state of the 4-queens problem. Consider neighbour relation B: swap the row positions of two queens. Which of the following is correct?

		Q	
			Q
	Q		
Q			

- (A) This state is a local optimum and is a global optimum.
- (B) This state is a local optimum and is NOT a global optimum.
- (C) This state is NOT a local optimum and is a global optimum.
- (D) This state is NOT a local optimum and NOT a global optimum.

CQ: Local and global optimum (2)

CQ: Consider the following state of the 4-queens problem. Consider neighbour relation A: move a single queen to another square in the same column. Which of the following is correct?

		Q	
			Q
	Q		
Q			

- (A) This state is a local optimum and is a global optimum.
- (B) This state is a local optimum and is NOT a global optimum.
- (C) This state is NOT a local optimum and is a global optimum.
- (D) This state is NOT a local optimum and NOT a global optimum.

Escaping a shoulder

- ▶ Sideway moves: allow the algorithm to move to a neighbour that has the same cost.
- ▶ Tabu list: keep a small list of recently visited states and forbid the algorithm to return to those states

Performance of hill climbing

- ▶ Perform quite well in practice.
- ▶ Makes rapid progress towards a solution.
Easy to improve a bad state.

8-queens problem: ≈ 17 million states.

- ▶ Basic hill climbing

% of instances solved: 14%

of steps until success/failure: 3-4 steps on average until success or failure.

- ▶ Basic hill climbing + ≤ 100 consecutive sideways moves:

% of instances solved: 94%

of steps until success/failure: 21 steps until success and 64 steps until failure.

Choosing the Neighbour Relation

How do we choose the neighbour relation?

- ▶ Small incremental change to the variable assignment

There's a trade-off:

- ▶ bigger neighbourhoods:
- ▶ smaller neighbourhoods:

Dealing with local optima

Hill climbing can get stuck at a local optimum.
What can we do?

- ▶ Restart search in a different part of the state space.
Hill climbing with random restarts
- ▶ Move to a state with a higher cost occasionally.
Simulated annealing

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Hill climbing

- Hill climbing with random restarts

- Simulated Annealing

- Genetic Algorithms

Revisiting the Learning goals

Hill climbing with random restarts

If at first you don't succeed, try, try again.

Restart the search with a randomly generated initial state when

- ▶ we found a local optimum, and
- ▶ we've found a plateau and made too many consecutive sideways moves.
- ▶ we've made too many moves.

Choose the best solution out of all the local optima found.

Will hill climbing + random restarts find the global optimum?

Will hill climbing with random restarts find the global optimum given enough time?

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Hill climbing

- Hill climbing with random restarts

- Simulated Annealing**

- Genetic Algorithms

Revisiting the Learning goals

Simulated Annealing

- ▶ Where do we start?

Start with a random solution and a large T .

- ▶ Which neighbour do we move to?

Choose a random neighbour.

If the neighbour is better than current, move to the neighbour.

If the neighbour is not better than the current state,
move to the neighbour with probability $p = e^{\Delta E/T}$.

- ▶ When do we stop?

Stop when $T = 0$.

Simulated Annealing

Algorithm 1 Simulated Annealing

```
1: current  $\leftarrow$  initial-state
2:  $T \leftarrow$  a large positive value
3: while  $T > 0$  do
4:   next  $\leftarrow$  a random neighbour of current
5:    $\Delta E \leftarrow$  current.cost - next.cost
6:   if  $\Delta E > 0$  then
7:     current  $\leftarrow$  next
8:   else
9:     current  $\leftarrow$  next with probability  $p = e^{\Delta E/T}$ 
10:  end if
11:  decrease  $T$ 
12: end while
13: return current
```

CQ: How does T affect $p = e^{\Delta E/T}$?

CQ: Consider a neighbour with a higher cost than the current node ($\Delta E < 0$).

As T decreases, how does $p = e^{\Delta E/T}$ change?
($p = e^{\Delta E/T}$ is the probability of moving to this neighbour.)

(A) As T decreases, $p = e^{\Delta E/T}$ increases.

(B) As T decreases, $p = e^{\Delta E/T}$ decreases.

CQ: How does ΔE affect $p = e^{\Delta E/T}$?

CQ: Assume that T is fixed. Consider a neighbour where $\Delta E < 0$

As ΔE decreases (becomes more negative),
how does $p = e^{\Delta E/T}$ change?

($p = e^{\Delta E/T}$ is the probability of moving to this neighbour.)

(A) As ΔE decreases, $p = e^{\Delta E/T}$ increases.

(B) As ΔE decreases, $p = e^{\Delta E/T}$ decreases.

Annealing Schedule

How should we decrease T ?

- ▶ Linear
- ▶ Logarithmic
- ▶ Exponential

If the temperature decreases slowly enough, simulated annealing is guaranteed to find the global optimum with probability approaching 1.

Examples of Simulated Annealing

- ▶ Example: getting a tennis ball into the deepest hole.
- ▶ Exploration versus exploitation

Learning Goals

Introduction to Local Search

Local Search Algorithms

- Hill climbing

- Hill climbing with random restarts

- Simulated Annealing

- Genetic Algorithms

Revisiting the Learning goals

Parallel Search

- ▶ **Idea:** maintain k nodes instead of one.
- ▶ At every stage, update each node.
- ▶ Whenever one node is a solution, report it.
- ▶ Like k restarts, but uses k times the minimum number of steps.

There's not really any reason to use this method. Why not?

Beam Search

- ▶ Maintain k nodes instead of one.
- ▶ Choose the k best nodes out of **all of the neighbors**.
- ▶ When $k = 1$, it is hill climbing.
- ▶ The value of k lets us limit space and parallelism.

Do you see any potential problem with beam search?

Stochastic Beam Search

- ▶ Choose the k nodes **probabilistically**.
- ▶ The probability that a neighbor is chosen is proportional to the fitness of the neighbour.
- ▶ Maintains diversity amongst the nodes.
- ▶ Asexual reproduction: each node mutates and the fittest offsprings survive.

Genetic algorithm

- ▶ Like stochastic beam search, but with sexual reproduction
Pairs of nodes are combined to create an offspring.
- ▶ Keep track of a set of states. Each state has a fitness.
- ▶ Randomly choose two states to reproduce.
The fitter a state, the most likely it's chosen to reproduce.
- ▶ Two parent states crossover to produce a child state.
- ▶ The child state mutates with a small independent probability.
- ▶ Add the child state to the new population. Repeat the steps above until we produce a new population. Replace the old population with the new one.
- ▶ Repeat until one state in the population has high enough fitness.

Genetic Algorithm

Algorithm 2 Genetic Algorithm

```
1:  $i = 0$ 
2: create initial population  $pop(i) = \{X_1, \dots, X_n\}$ 
3: while true do
4:   if  $\exists x \in pop(i)$  with high enough  $f(x)$  then
5:     break
6:   end if
7:   for each  $X_i \in pop(i)$  calculate  $pr(X_i) = f(X_i) / \sum_i f(X_i)$ 
8:   for  $j$  from 1 to  $n$  do
9:     choose  $a$  randomly based on  $pr(X_i)$ 
10:    choose  $b$  randomly based on  $pr(X_i)$ 
11:    child  $\leftarrow$  crossover( $a, b$ )
12:    child mutates with small probability
13:    add child to  $pop(i+1)$ 
14:   end for
15:    $i = i + 1$ 
16: end while
17: return  $x \in pop(i)$  with highest fitness
```

Comparing hill climbing and genetic algorithm

- ▶ How does each algorithm explore the state space?

Hill climbing generates neighbours of the state based on the neighbour relation.

Genetic algorithm ...

- ▶ How does each algorithm optimize the quality of the state/population?

Hill climbing moves to the best neighbour.

Genetic algorithm ...

Revisiting the Learning Goals

By the end of the lecture, you should be able to

- ▶ Describe the advantages of local search over other search algorithms.
- ▶ Formulate a real world problem as a local search problem.
- ▶ Given a local search problem, verify whether a state is a local/global optimum.
- ▶ Describe strategies for escaping local optima.
- ▶ Trace the execution of hill climbing, hill climbing with random restarts, and simulated annealing.
- ▶ Compare and contrast the properties of local search algorithms.