The Best Medium-Hard Data Analyst SQL Interview Questions

By Zachary Thomas (zthomas.nc@gmail.com, Twitter, LinkedIn)

Tip: See the Table of Contents (document outline) by hovering over the vertical line on the right side of the page

Update: Thanks everyone for the support and feedback! See this discussion on this post on Hacker News, Linkedin, Eric Weber's Linkedin post

Background & Motivation

> The first 70% of SQL is pretty straightforward but the remaining 30% can be pretty tricky.

Between the fall of 2015 and the summer of 2019 I interviewed for data analyst and data scientists positions four separate times, getting to onsite interviews at over a dozen companies. After an interview in 2017 went poorly – mostly due to me floundering at the more difficult SQL questions they asked me – I started putting together a study guide of medium and hard SQL questions to better prepare and found it particularly useful during my 2019 interview cycle. Over the past year I have shared that guide with a couple of friends, and with the extra time on my hands due to the coronavirus pandemic, I have polished it up into this doc.

There are plenty of great beginner SQL guides out there. My favorites are Codecademy's interactive SQL courses and Zi Chong Kao's Select Star SQL. However, like I told a friend, while the first 70% of SQL is pretty straightforward, the remaining 30% can be pretty tricky. Data analyst and data scientist interview questions at technology companies often pull from that 30%.

Strangely, I have never really found a comprehensive source online for those medium-hard SQL questions, which is why I put together this guide.

Working through this guide should improve your performance on data analyst interviews. It should also make you better at your current and future job positions. Personally, I find some of the SQL patterns found in this doc useful for ETLs powering reporting tools featuring trends over time.

To be clear, data analyst and data scientist interviews consist of more than SQL questions. Other common topics include explaining past projects, A/B testing (I like Udacity's course on the subject), metric development and open-ended analytical problems. This Quora answer has Facebook's product analyst interview guide circa 2017, which discusses this topic in more depth. That said, if improving your SQL skills can make your interviews less stressful than they already are, it could very well be worth your time.

In the future, I may transition this doc to a website like Select Star SQL with an embedded SQL editor so that readers can write SQL statements to questions and get real-time feedback on their code. Another option could be adding these questions as problems on Leetcode. For the time being though I just wanted to publish this doc so that people could find it useful now.

I would love to get your feedback on this doc. Please drop a note if you find this useful, have improvements/corrections, or encounter other good resources for medium/hard difficulty SQL questions.

Assumptions & How to use this guide

Assumptions about SQL proficiency: This guide assumes you have a working knowledge of SQL. You probably use it frequently at work already but want to sharpen your skills on topics like self-joins and window functions.

How to use this guide: Since interviews usually utilize a whiteboard or a virtual (non-compiling) notepad, my recommendation is to get out a pencil and paper and write out your solutions to each part of the problem, and once complete compare your answers to the answer key. Or, complete these with a friend who can act as the interviewer!

- Small SQL syntax errors aren't a big deal during whiteboard/notepad interviews. However, they can be distracting to the interviewer, so ideally practice reducing these so your logic shines through in the interview.
- The answers I provide may not be the only way to successfully solve the question. Feel free to message with additional solutions and I can add them to this guide!

Tips on solving difficult SQL interview questions

This advice mirrors typical code interview advice ...

1. Listen carefully to problem description, repeat back the crux of the problem to the interviewer
2. Spell out an edge case to demonstrate you actually understand problem (i.e. a row that *wouldn't* be included in the output of the SQL you are about to sketch out)
3. (If the problem involves a self-join) For your own benefit sketch out what the self-join will look like – this will typically be at least three columns: a column of interest from the main table, the column to join from the main table, and the column to join from the secondary table
   a. Or, as you get more used to self-join problems, you can explain this step verbally
4. Start writing SQL – err towards writing SQL versus trying to perfectly understand the problem. Verbalize your assumptions as you go so your interviewer can correct you if you go astray.

Acknowledgments and Additional Resources

Some of the problems listed here are adapted from old Periscope blog posts (mostly written around 2014 by Sean Cook, although his authorship seems to have been removed from the posts following SiSense's merger with Periscope) or discussions from Stack Overflow; I've noted them at the start of questions as appropriate.

Select Star SQL has good challenge questions that are complementary to the questions in this doc.

Please note that these questions are not literal copies of SQL interview questions I have encountered while interviewing nor were they interview questions used at a company I have worked at or work at.

Self-Join Practice Problems

#1: MoM Percent Change

Context: Oftentimes it's useful to know how much a key metric, such as monthly active users, changes between months. Say we have a table `logins` in the form:

```
| user_id | date       |
|---------|------------|
| 1       | 2018-07-01 |
| 234     | 2018-07-02 |
```

```
| 3         | 2018-07-02 |
| 1         | 2018-07-02 |
| ...       | ...        |
| 234       | 2018-10-04 |
```

Task: Find the month-over-month percentage change for monthly active users (MAU).

Solution:

(This solution, like other solution code blocks you will see in this doc, contains comments about SQL syntax that may differ between flavors of SQL or other comments about the solutions as listed)

```sql
WITH mau AS
(
  SELECT
    /*
     * Typically, interviewers allow you to write psuedocode for date functions
     * i.e. will NOT be checking if you have memorized date functions.
     * Just explain what your function does as you whiteboard
     *
     * DATE_TRUNC() is available in Postgres, but other SQL date functions or
     * combinations of date functions can give you a identical results
     * See https://www.postgresql.org/docs/9.0/functions-datetime.html#FUNCTIONS-DATETI
     */
    DATE_TRUNC('month', date) month_timestamp,
    COUNT(DISTINCT user_id) mau
  FROM
    logins
  GROUP BY
    DATE_TRUNC('month', date)
 )

  SELECT
    /*
     * You don't literally need to include the previous month in this SELECT statement.
     *
     * However, as mentioned in the "Tips" section of this guide, it can be helpful
     * to at least sketch out self-joins to avoid getting confused which table
     * represents the prior month vs current month, etc.
     */
    a.month_timestamp previous_month,
    a.mau previous_mau,
    b.month_timestamp current_month,
    b.mau current_mau,
    ROUND(100.0*(b.mau - a.mau)/a.mau,2) AS percent_change
  FROM
    mau a
  JOIN
    /*
     * Could also have done `ON b.month_timestamp = a.month_timestamp + interval '1 mor
     */
```

```
        mau b ON a.month_timestamp = b.month_timestamp - interval '1 month'
```

#2: Tree Structure Labeling

Context: Say you have a table `tree` with a column of nodes and a column corresponding parent nodes

```
node    parent
1         2
2         5
3         5
4         3
5         NULL
```

Task: Write SQL such that we label each node as a "leaf", "inner" or "Root" node, such that for the nodes above we get:

```
node    label
1        Leaf
2        Inner
3        Inner
4        Leaf
5        Root
```

A solution which works for the above example will receive full credit, although you can receive extra credit for providing a solution that is generalizable to a tree of any depth (not just depth = 2, as is the case in the example above).

(Side note: this link has more details on Tree data structure terminology. Not needed to solve the problem though!)

Solution:

Note: This solution works for the example above with tree depth = 2, but is not generalizable beyond that.

```
WITH join_table AS
(
SELECT
    a.node a_node,
    a.parent a_parent,
    b.node b_node,
    b.parent b_parent
 FROM
    tree a
 LEFT JOIN
    tree b ON a.parent = b.node
 )

 SELECT
    a_node node,
```

```
      CASE
          WHEN b_node IS NULL and b_parent IS NULL THEN 'Root'
          WHEN b_node IS NOT NULL and b_parent IS NOT NULL THEN 'Leaf'
          ELSE 'Inner'
      END AS label
  FROM
      join_table
```

An alternate solution, that is generalizable to any tree depth:

Acknowledgement: this more generalizable solution was contributed by Fabian Hofmann on 5/2/20. Thank, FH!

```
WITH join_table AS
(
    SELECT
        cur.node,
        cur.parent,
        COUNT(next.node) AS num_children
    FROM
        tree cur
    LEFT JOIN
        tree next ON (next.parent = cur.node)
    GROUP BY
        cur.node,
        cur.parent
)

SELECT
    node,
    CASE
        WHEN parent IS NULL THEN "Root"
        WHEN num_children = 0 THEN "Leaf"
        ELSE "Inner"
    END AS label
FROM
    join_table
```

An alternate solution, without explicit joins:

Acknowledgement: William Chargin on 5/2/20 noted that `WHERE parent IS NOT NULL` is needed to make this solution return `Leaf` instead of `NULL`. Thanks, WC!

```
SELECT
    node,
    CASE
        WHEN parent IS NULL THEN 'Root'
        WHEN node NOT IN
            (SELECT parent FROM tree WHERE parent IS NOT NULL) THEN 'Leaf'
        WHEN node IN (SELECT parent FROM tree) AND parent IS NOT NULL THEN 'Inner'
    END AS label
from
    tree
```

#3: Retained Users Per Month (multi-part)

Acknowledgement: this problem is adapted from SiSense's "Using Self Joins to Calculate Your Retention, Churn, and Reactivation Metrics" blog post

PART 1:

Context: Say we have login data in the table `logins`:

```
| user_id | date       |
|---------|------------|
| 1       | 2018-07-01 |
| 234     | 2018-07-02 |
| 3       | 2018-07-02 |
| 1       | 2018-07-02 |
| ...     | ...        |
| 234     | 2018-10-04 |
```

Task: Write a query that gets the number of retained users per month. In this case, retention for a given month is defined as the number of users who logged in that month who also logged in the immediately previous month.

Solution:

```sql
SELECT
    DATE_TRUNC('month', a.date) month_timestamp,
    COUNT(DISTINCT a.user_id) retained_users
 FROM
    logins a
 JOIN
    logins b ON a.user_id = b.user_id
        AND DATE_TRUNC('month', a.date) = DATE_TRUNC('month', b.date) +
                                          interval '1 month'
 GROUP BY
    date_trunc('month', a.date)
```

Acknowledgement: Tom Moertel pointed out de-duping user-login pairs before the self-join would make the solution more efficient and contributed the alternate solution below. Thanks, TM!

Note: De-duping `logins` would also make the given solutions to Parts 2 and 3 of this problem more efficient as well.

Alternate solution:

```sql
WITH DistinctMonthlyUsers AS (
  /*
  * For each month, compute the *set* of users having logins.
  */
    SELECT DISTINCT
      DATE_TRUNC('MONTH', date) AS month_timestamp,
```

```
      user_id
    FROM logins
  )

SELECT
  CurrentMonth.month_timestamp month_timestamp,
  COUNT(PriorMonth.user_id) AS retained_user_count
FROM
    DistinctMonthlyUsers AS CurrentMonth
LEFT JOIN
    DistinctMonthlyUsers AS PriorMonth
  ON
    CurrentMonth.month_timestamp = PriorMonth.month_timestamp + INTERVAL '1 MONTH'
    AND
    CurrentMonth.user_id = PriorMonth.user_id
```

PART 2:

Task: Now we'll take retention and turn it on its head: Write a query to find many users last month *did not* come back this month. i.e. the number of churned users.

Solution:

```
SELECT
    DATE_TRUNC('month', a.date) month_timestamp,
    COUNT(DISTINCT b.user_id) churned_users
FROM
    logins a
FULL OUTER JOIN
    logins b ON a.user_id = b.user_id
        AND DATE_TRUNC('month', a.date) = DATE_TRUNC('month', b.date) +
                                          interval '1 month'
WHERE
    a.user_id IS NULL
GROUP BY
    DATE_TRUNC('month', a.date)
```

Note that there are solutions to this problem that can use `LEFT` or `RIGHT` joins.

PART 3:

Context: You now want to see the number of active users this month *who have been reactivated* – in other words, users who have churned but this month they became active again. Keep in mind a user can reactivate after churning *before* the previous month. An example of this could be a user active in February (appears in `logins`), no activity in March and April, but then active again in May (appears in `logins`), so they count as a reactivated user for May .

Task: Create a table that contains the number of reactivated users per month.

Solution:

```
    SELECT
        DATE_TRUNC('month', a.date) month_timestamp,
        COUNT(DISTINCT a.user_id) reactivated_users,
        /*
        * At least in the flavors of SQL I have used, you don't need to
        * include the columns used in HAVING in the SELECT statement.
        * I have written them out for clarity here.
        */
        MAX(DATE_TRUNC('month', b.date)) most_recent_active_previously
    FROM
        logins a
    JOIN
        logins b ON a.user_id = b.user_id
                AND
                DATE_TRUNC('month', a.date) > DATE_TRUNC('month', b.date)
    GROUP BY
        DATE_TRUNC('month', a.date)
    HAVING
        month_timestamp > most_recent_active_previously + interval '1 month'
```

#4: Cumulative Sums

Acknowledgement: This problem was inspired by Sisense's "Cash Flow modeling in SQL" blog post

Context: Say we have a table `transactions` in the form:

```
| date       | cash_flow |
|------------|-----------|
| 2018-01-01 | -1000     |
| 2018-01-02 | -100      |
| 2018-01-03 | 50        |
| ...        | ...       |
```

Where `cash_flow` is the revenues minus costs for each day.

Task: Write a query to get *cumulative* cash flow for each day such that we end up with a table in the form below:

```
| date       | cumulative_cf |
|------------|---------------|
| 2018-01-01 | -1000         |
| 2018-01-02 | -1100         |
| 2018-01-03 | -1050         |
| ...        | ...           |
```

Solution:

```
SELECT
    a.date date,
```

```
    SUM(b.cash_flow) as cumulative_cf
FROM
    transactions a
JOIN b
    transactions b ON a.date >= b.date
GROUP BY
    a.date
ORDER BY
    date ASC
```

Alternate solution using a window function (more efficient!):

```
SELECT
    date,
    SUM(cash_flow) OVER (ORDER BY date ASC) as cumulative_cf
FROM
    transactions
ORDER BY
    date ASC
```

#5: Rolling Averages

Acknowledgement: This problem is adapted from Sisense's "Rolling Averages in MySQL and SQL Server" blog post

Note: there are different ways to compute rolling/moving averages. Here we'll use a preceding average which means that the metric for the 7th day of the month would be the average of the preceding 6 days and that day itself.

Context: Say we have table `signups` in the form:

```
| date       | sign_ups |
|------------|----------|
| 2018-01-01 | 10       |
| 2018-01-02 | 20       |
| 2018-01-03 | 50       |
| ...        | ...      |
| 2018-10-01 | 35       |
```

Task: Write a query to get 7-day rolling (preceding) average of daily sign ups.

Solution:

```
SELECT
  a.date,
  AVG(b.sign_ups) average_sign_ups
FROM
  signups a
JOIN
  signups b ON a.date <= b.date + interval '6 days' AND a.date >= b.date
GROUP BY
  a.date
```

#6: Multiple Join Conditions

Acknowledgement: This problem was inspired by Sisense's "Analyzing Your Email with SQL" blog post

Context: Say we have a table `emails` that includes emails sent to and from `zach@g.com`:

```
| id | subject  | from           | to             | timestamp           |
|----|----------|----------------|----------------|---------------------|
| 1  | Yosemite | zach@g.com     | thomas@g.com   | 2018-01-02 12:45:03 |
| 2  | Big Sur  | sarah@g.com    | thomas@g.com   | 2018-01-02 16:30:01 |
| 3  | Yosemite | thomas@g.com   | zach@g.com     | 2018-01-02 16:35:04 |
| 4  | Running  | jill@g.com     | zach@g.com     | 2018-01-03 08:12:45 |
| 5  | Yosemite | zach@g.com     | thomas@g.com   | 2018-01-03 14:02:01 |
| 6  | Yosemite | thomas@g.com   | zach@g.com     | 2018-01-03 15:01:05 |
| .. | ..       | ..             | ..             | ..                  |
```

Task: Write a query to get the response time per email (`id`) sent to `zach@g.com`. Do not include `id`s that did not receive a response from `zach@g.com`. Assume each email thread has a unique subject. Keep in mind a thread may have multiple responses back-and-forth between `zach@g.com` and another email address.

Solution:

```
SELECT
    a.id,
    MIN(b.timestamp) - a.timestamp as time_to_respond
FROM
    emails a
JOIN
    emails b
        ON
            b.subject = a.subject
        AND
            a.to = b.from
        AND
            a.from = b.to
        AND
            a.timestamp < b.timestamp
  WHERE
    a.to = 'zach@g.com'
  GROUP BY
    a.id
```

Window Function Practice Problems

#1: Get the ID with the highest value

Context: Say we have a table `salaries` with data on employee salary and department in the following format:

```
  depname   | empno | salary |
------------+-------+--------+
 develop    |    11 |   5200 |
 develop    |     7 |   4200 |
 develop    |     9 |   4500 |
 develop    |     8 |   6000 |
 develop    |    10 |   5200 |
 personnel  |     5 |   3500 |
 personnel  |     2 |   3900 |
 sales      |     3 |   4800 |
 sales      |     1 |   5000 |
 sales      |     4 |   4800 |
```

Task: Write a query to get the `empno` with the highest salary. Make sure your solution can handle ties!

Solution:

```
WITH max_salary AS (
    SELECT
        MAX(salary) max_salary
    FROM
        salaries
    )
SELECT
    s.empno
FROM
    salaries s
JOIN
    max_salary ms ON s.salary = ms.max_salary
```

Alternate solution using `RANK()`:

```
WITH sal_rank AS
  (SELECT
     empno,
     RANK() OVER(ORDER BY salary DESC) rnk
   FROM
     salaries)
SELECT
  empno
FROM
  sal_rank
WHERE
  rnk = 1;
```

#2: Average and rank with a window function (multi-part)

PART 1:

Context: Say we have a table `salaries` in the format:

```
   depname  | empno | salary |
------------+-------+--------+
  develop   |    11 |   5200 |
  develop   |     7 |   4200 |
  develop   |     9 |   4500 |
  develop   |     8 |   6000 |
  develop   |    10 |   5200 |
  personnel |     5 |   3500 |
  personnel |     2 |   3900 |
  sales     |     3 |   4800 |
  sales     |     1 |   5000 |
  sales     |     4 |   4800 |
```

Task: Write a query that returns the same table, but with a new column that has average salary per `depname`. We would expect a table in the form:

```
   depname  | empno | salary | avg_salary |
------------+-------+--------+------------+
  develop   |    11 |   5200 |       5020 |
  develop   |     7 |   4200 |       5020 |
  develop   |     9 |   4500 |       5020 |
  develop   |     8 |   6000 |       5020 |
  develop   |    10 |   5200 |       5020 |
  personnel |     5 |   3500 |       3700 |
  personnel |     2 |   3900 |       3700 |
  sales     |     3 |   4800 |       4867 |
  sales     |     1 |   5000 |       4867 |
  sales     |     4 |   4800 |       4867 |
```

Solution:

```
SELECT
    *,
    /*
    * AVG() is a Postgres command, but other SQL flavors like BigQuery use
    * AVERAGE()
    */
    ROUND(AVG(salary),0) OVER (PARTITION BY depname) avg_salary
FROM
    salaries
```

PART 2:

Task: Write a query that adds a column with the rank of each employee based on their salary within their department, where the employee with the highest salary gets the rank of `1`. We would expect a table in the form:

```
   depname   | empno | salary | salary_rank |
-------------+-------+--------+-------------+
   develop   |    11 |   5200 |           2 |
   develop   |     7 |   4200 |           5 |
   develop   |     9 |   4500 |           4 |
   develop   |     8 |   6000 |           1 |
   develop   |    10 |   5200 |           2 |
   personnel |     5 |   3500 |           2 |
   personnel |     2 |   3900 |           1 |
   sales     |     3 |   4800 |           2 |
   sales     |     1 |   5000 |           1 |
   sales     |     4 |   4800 |           2 |
```

Solution:

```
SELECT
    *,
    RANK() OVER(PARTITION BY depname ORDER BY salary DESC) salary_rank
FROM
    salaries
```

Other Medium/Hard SQL Practice Problems

#1: Histograms

Context: Say we have a table `sessions` where each row is a video streaming session with length in seconds:

```
| session_id | length_seconds |
|------------|----------------|
| 1          | 23             |
| 2          | 453            |
| 3          | 27             |
| ..         | ..             |
```

Task: Write a query to count the number of sessions that fall into bands of size 5, i.e. for the above snippet, produce something akin to:

```
| bucket   | count |
|----------|-------|
| 20-25    | 2     |
| 450-455  | 1     |
```

Get complete credit for the proper string labels ("5-10", etc.) but near complete credit for something that is communicable as the bin.

Solution:

```
WITH bin_label AS
(SELECT
    session_id,
    FLOOR(length_seconds/5) as bin_label
 FROM
    sessions
 )
 SELECT
    CONCATENTATE(STR(bin_label*5), '-', STR(bin_label*5+5)) bucket,
    COUNT(DISTINCT session_id) count
 GROUP BY
    bin_label
 ORDER BY
    bin_label ASC
```

#2: CROSS JOIN (multi-part)

PART 1:

Context: Say we have a table `state_streams` where each row is a state and the total number of hours of streaming from a video hosting service:

```
| state | total_streams |
|-------|---------------|
| NC    | 34569         |
| SC    | 33999         |
| CA    | 98324         |
| MA    | 19345         |
| ..    | ..            |
```

(In reality these kinds of aggregate tables would normally have a date column, but we'll exclude that component in this problem)

Task: Write a query to get the pairs of states with total streaming amounts within 1000 of each other. For the snippet above, we would want to see something like:

```
| state_a | state_b |
|---------|---------|
| NC      | SC      |
| SC      | NC      |
```

Solution:

```
SELECT
    a.state as state_a,
    b.state as state_b
```

```
    FROM
        state_streams a
    CROSS JOIN
        state_streams b
    WHERE
        ABS(a.total_streams - b.total_streams) < 1000
        AND
        a.state <> b.state
```

FYI, `CROSS JOIN`s can also be written without explicitly specifying a join:

```
SELECT
    a.state as state_a,
    b.state as state_b
 FROM
    state_streams a, state_streams b
 WHERE
    ABS(a.total_streams - b.total_streams) < 1000
    AND
    a.state <> b.state
```

PART 2:

Note: This question is considered more of a bonus problem than an actual SQL pattern. Feel free to skip it!

Task: How could you modify the SQL from the solution to Part 1 of this question so that duplicates are removed? For example, if we used the sample table from Part 1, the pair NC and SC should only appear in one row instead of two.

Solution:

```
SELECT
    a.state as state_a,
    b.state as state_b
 FROM
    state_streams a, state_streams b
 WHERE
    ABS(a.total_streams - b.total_streams) < 1000
    AND
    a.state > b.state
```

#3: Advancing Counting

Acknowledgement: This question is adapted from this Stack Overflow question by me (zthomas.nc)

Note: this question is probably more complex than the kind you would encounter in an interview. Consider it a challenge problem, or feel free to skip it!

Context: Say I have a table `table` in the following form, where a `user` can be mapped to multiple values of `class`:

```
| user | class |
|------|-------|
| 1    | a     |
| 1    | b     |
| 1    | b     |
| 2    | b     |
| 3    | a     |
```

Task: Assume there are only two possible values for `class`. Write a query to count the number of users in each class such that any user who has label `a` and `b` gets sorted into `b`, any user with just `a` gets sorted into `a` and any user with just `b` gets into `b`.

For `table` that would result in the following table:

```
| class | count |
|-------|-------|
| a     | 1     |
| b     | 2     |
```

Solution:

```sql
WITH usr_b_sum AS
(
    SELECT
        user,
        SUM(CASE WHEN class = 'b' THEN 1 ELSE 0 END) num_b
    FROM
        table
    GROUP BY
        user
),

usr_class_label AS
(
    SELECT
        user,
        CASE WHEN num_b > 0 THEN 'b' ELSE 'a' END class
    FROM
        usr_b_sum
)

SELECT
    class,
    COUNT(DISTINCT user) count
FROM
    usr_class_label
GROUP BY
    class
```

```
ORDER BY
    class ASC
```

Alternate solution: Using SELECTs in the SELECT statement and UNION:

```
SELECT
    "a" class,
    COUNT(DISTINCT user_id) -
        (SELECT COUNT(DISTINCT user_id) FROM table WHERE class = 'b') count
UNION
SELECT
    "b" class,
    (SELECT COUNT(DISTINCT user_id) FROM table WHERE class = 'b') count
```

Alternate solution: Since the problem as stated didn't ask for generalizable solution, you can leverage the fact that b > a to produce this straightforward solution:

Acknowledgement: Thanks to Karan Gadiya for contributing this solution. Thanks, KG!

```
WITH max_class AS (
    SELECT
        user,
        MAX(class) as class
    FROM
        table
    GROUP BY
        user
)

SELECT
    class,
    COUNT(user)
FROM
    max_class
GROUP BY
    class
```